

---

# Numerical Methods

## Lecture 7

---

Nicholas Christian  
BIOST 2094 Spring 2011

# Outline

---

1. Integration
2. Differentiation
3. Root Finding
4. Optimization
5. Adaptive Rejection Sampling

# Numerical Integration

---

- The function `integrate()` is used to integrate functions of one variable over a finite or infinite interval
- When integrating over infinite intervals use `-Inf` or `Inf`
- The function you are integrating must accept a vector argument and return a vector with the function evaluated at each point in the argument vector. Use `Vectorize()` to create a function that performs vector computations.
- Using `integrate()` along with `sapply()` allows you to evaluate multiple integrals.

```
integrate(f, lower, upper, ...)
```

`f` Function integrating

`lower` Lower limit of integration

`upper` Upper limit of integration

`...` Additional arguments passed to `f`

There are additional arguments for controlling the accuracy of the estimate

## Example - Numerical Integration

---

```
# Simple example
f <- function(x) exp(x)

# Returns a class "integrate" object
(I <- integrate(f, lower=0, upper=1))
names(I)

# Get the estimate of the intergral
integrate(f, lower=0, upper=1)$value

# Integrate a normal density function
integrate(dnorm, -1.96, 1.96, mean=0, sd=1)
integrate(dnorm, -Inf, Inf, mean=0, sd=1)

# Vector Computations
g <- function(x) 1
integrate(g, 0, 1)

# Need to vectorize g
g(1:10)
Vectorize(g)(1:10)
integrate(Vectorize(g), 1, 10)
```

## Example - Numerical Multiple Integration

---

- The following trick came from the R message board

$$\int_0^3 \int_1^2 x^2 y \, dy dx$$

```
# Iterated Integral
integrate(function(x) {
  sapply(x, function(x) {
    integrate(function(y) x^2*y, 1, 2)$value
  })
}, 0, 3)
```

$$\int_0^1 \int_x^1 x \sin(y^2) \, dy dx$$

```
# Double Integral
integrate(function(x) {
  sapply(x, function(x) {
    integrate(function(y) x*sin(y^2), x, 1)$value
  })
}, 0, 1)
```

# Symbolic Derivatives

---

- To perform symbolic derivatives of simple expressions use `deriv()`

```
deriv(expr, namevec, hessian=FALSE)
```

<code>expr</code>	Either a formula with no left-hand side or an expression
<code>namevec</code>	Character vector, giving the variable names with respect to which derivatives will be computed
<code>hessian</code>	Logical whether or not the second derivatives should be computed and returned

- Returns an unevaluated call for computing the `expr` along with a `gradient` attribute containing the gradient matrix and if `hessian=TRUE` a `hessian` attribute containing the Hessian array

## Example - Symbolic Derivatives

---

```
# Formula argument
dx.expr <- deriv(~x^2, "x", hessian=TRUE) # Unevaluated expression
dx.expr

x <- 1:5
dx.num <- eval(dx.expr) # Evaluated expression
dx.num
attr(dx.num, "gradient") # Gradient Vector

# Expression argument
dx.dy.expr <- deriv(expression(x^2*y^4), c("x", "y"), hessian=TRUE)
dx.dy.expr

x <- 1:5; y <- 1:5
dx.dy.num <- eval(dx.dy.expr)
dx.dy.num
attr(dx.dy.num, "hessian") # For matrix "x", the columns are
                           # d2/dx2 and d2/dxdy
```

# Root Finding

---

- The function `uniroot()` searches an interval for the root of a function

`uniroot(f, interval, ...)`

<code>f</code>	Function to get the root of
<code>interval</code>	Vector giving the interval <code>c(lower, upper)</code> to be searched
<code>...</code>	Additional arguments to be passed to <code>f</code>

- *Only returns one root*, does not return multiple roots, if multiple roots exist



## Example - Root Finding

---

```
# Function we want to find the roots of
f <- function(x) {-x^4-.5*x^3+9*x^2-x-5}

# Plot the function we want to find the roots of
curve(f(x), -3, 3, lwd=2, main=expression(f(x)=-x^4-.5*x^3+9*x^2-x-5))
abline(h=0, lty=2, lwd=1)
points(c(2.563, .866, -.697), c(0,0,0), pch=19, cex=1.5, col="blue")

# Find the root between 0 and 2
uniroot(f, c(0, 2))

# Only returns one root, even if there are multiple roots in the interval
uniroot(f, c(-3, 3))

# The f values at the end points need to be of opposite sign
uniroot(f, c(1, 2))
```

## Example - Simulating Survival Data

- Suppose we want to simulate  $n = 100$  observations. The event times  $T$  follow a Weibull distribution with shape parameter  $a = 2$  and scale parameter  $b = 1$  and the censoring times  $C$  are distributed uniformly from 0 to  $\tau$ . Then the observed time is,  $X = \min(T, C)$ . Assume  $T$  and  $C$  are independent.
- The value of  $\tau$  is chosen to achieve the desired censoring rate. For this example we want 25% of the observations to be censored. Thus,  $P(C < T) = 0.25$  and we need to choose  $\tau$  such that,

$$\begin{aligned} P(C < T) &= \int_0^\tau \int_c^\infty f_T(t) f_C(c) dt dc \\ &= \frac{1}{\tau} \int_0^\tau \int_c^\infty f_T(t) dt dc \\ &= \frac{1}{\tau} \int_0^\tau S_T(c) dc \\ &= 0.25 \end{aligned}$$

## Example - Simulating Survival Data

- Find  $\tau$  by using `integrate()` and `uniroot()` to solve,

$$\frac{1}{\tau} \int_0^{\tau} S_T(c) dt dc - 0.25 = 0$$

```
# Calculate tau for a uniform(0, tau) censoring distribution in order to get the
# desired censoring rate (pctCensor) when the event times are Weibull(a, b)
uniformCensorUpperBound <- function(pctCensor, a, b) {
  f <- function(tau) {
    S <- function(t) {1-pweibull(t, shape=a, scale=b)}
    (integrate(S, 0, tau)$value)/tau - pctCensor
  }
  return(uniroot(f, c(0.001, 100))$root)
}

(tau <- uniformCensorUpperBound(.25, 2, 1))

# Verify
T <- rweibull(100, shape=2, scale=1)      # Event times
C <- runif(100, 0, tau)                   # Censor times
index <- apply(cbind(C, T), 1, which.min)-1 # Event indicator
table(index)/100
```

# Optimization

---

- R has several functions for optimization,
 

optimize()	One dimensional optimization, no gradient or Hessian
optim()	General purpose optimization, five possible methods, gradient optional
constrOptim()	Minimization of a function subject to linear inequality constraints, gradient optional
nlm()	Non-linear minimization, can optionally include the gradient and hessian of the function as attributes of the objective function
nlminb()	Minimization using PORT routines, can optionally include the gradient and Hessian of the objective function as additional arguments
- These functions use different algorithms and accept different arguments, no one function is superior to the others. Which function you use depends on your particular problem; use `optimize()` for one-dimensional problems.
- To turn a minimization problem into a maximization problem, multiply the objective function and gradient by -1
- There are also packages with additional optimization functions, <http://cran.r-project.org/web/views/Optimization.html>

# One-Dimensional Optimization

---

```
optimize(f, interval, ..., maximum=FALSE)
```

f	Function to be optimized
interval	Vector giving the interval c(lower, upper) to be searched
...	Additional arguments to be passed to f
maximum	Logical, find maximum if TRUE

- Cannot specify the gradient to assist with the optimization

## Example - optimize()

---

```
# Objective function
f <- function(x) {-x^4-.5*x^3+9*x^2-x-5}

# Plot the objective function
curve(f(x), -3, 3, lwd=2, main=expression(f(x)=-x^4-.5*x^3+9*x^2-x-5))
abline(h=0, lty=2, lwd=1)

# Find the minimum of f
(min.f <- optimize(f, c(-3, 3)))
points(min.f$minimum, min.f$objective, pch=19, cex=1.5, col="blue")

# Find the maximum of f
(max.f <- optimize(f, c(-3, 3), maximum=TRUE))
points(max.f$maximum, max.f$objective, pch=19, cex=1.5, col="green3")

# Becareful, may not get a global max
(max.f <- optimize(f, c(1, 3), maximum=TRUE))
points(max.f$maximum, max.f$objective, pch=19, cex=1.5, col="red")
```

# Multi-Dimensional Optimization

---

```
optim(par, fn, gr=NULL, ..., method, control)
```

<code>par</code>	Initial values
<code>fn</code>	Function to be optimized, argument is a vector of parameters
<code>gr</code>	A function that returns the gradient, same argument as <code>fn</code>
<code>...</code>	Additional arguments passed to <code>fn</code>
<code>method</code>	Method to be used
<code>control</code>	List of control parameters (number of iterations, tolerance, etc.)

- By default the minimum is found, to find the maximum set the control parameter `fnscale` to `-1`, `control=list(fnscale=-1)`. This divides the objective function and gradient by `-1`.

## Example - optim()

---

```
# Two-dimensional objective function, the argument needs to be a vector
f <- function(x) {
  x1 = x[1]
  x2 = x[2]
  z = 10*x1^2*x2 - 5*x1^2 - 4*x2^2-x1^4-2*x2^4
  return(z)
}

# Plot the objective function
x <- y <- seq(-4, 4, len=50)
z <- outer(x, y, FUN=function(x,y) apply(cbind(x,y), 1, f))
filled.contour(x, y, z,
  color.palette=colorRampPalette(c("blue4","blue3","white","green3","green4")))

# Find the maximum points, from the contour plot there appears to be two maximums
# convergence=0 means convergence
# convergence=1 means reached max number of iterations
(pt1 <- optim(c(-2, 2), f, control=list(fnscale=-1)))
(pt2 <- optim(c(2, 2), f, control=list(fnscale=-1)))
```



## Example - optim()

---

```
# Plot maximum points on contour plot
filled.contour(x, y, z,
  color.palette=colorRampPalette(c("blue4","blue3","white","green3","green4")),
  plot.axes={
    axis(1); axis(2)
    points(pt1$par[1], pt1$par[2], pch=19, col="red", cex=1.5)
    text(pt1$par[1], pt1$par[2], round(pt1$value, 3), pos=1)
    points(pt2$par[1], pt2$par[2], pch=19, col="red", cex=1.5)
    text(pt2$par[1], pt2$par[2], round(pt2$value, 3), pos=1)
  })

# Find maximum points using gradient
g <- function(x) {
  x1 = x[1]
  x2 = x[2]
  dx1 <- 20*x1*x2 - 10*x1 - 4*x1^3
  dx2 <- 10*x1^2 - 8*x2 - 8*x2^3
  return(c(dx1, dx2))
}

# Not all methods use the gradient, default method "Nelder-Mead" does not
optim(c(-2, 2), fn=f, gr=g, method="BFGS", control=list(fnscale=-1))
```

# Constrained Optimization

---

- For box constraints use the "L-BFGS-B" method in `optim()` and the arguments `lower` and `upper` to give bounds for the arguments
- For linear inequality constraints use `constrOptim()`

```
constrOptim(theta, f, grad, ui, ci, control, ...)
```

<code>theta</code>	Starting values, $p \times 1$ vector
<code>f</code>	Function to be optimized
<code>grad</code>	Function that returns the gradient
<code>ui</code>	Constraint matrix, $k \times p$
<code>ci</code>	Constraint vector, $k \times 1$ vector
<code>control</code>	List of control parameters
<code>...</code>	Additional arguments passed to <code>f</code>

- Feasible region is defined by `ui %*% theta - ci >= 0`

## Example - Constrained Optimization

---

```
f <- function(x) {          # Objective function
  x1 = x[1]; x2 = x[2]
  10*x1^2*x2 - 5*x1^2 - 4*x2^2-x1^4-2*x2^4
}

g <- function(x) {          # Gradient
  x1 = x[1]; x2 = x[2]
  dx1 <- 20*x1*x2 - 10*x1 - 4*x1^3; dx2 <- 10*x1^2 - 8*x2 - 8*x2^3
  return(c(dx1, dx2))
}

# Box constraints, x1>0
optim(c(1,0), f, lower=c(0, -Inf), method="L-BFGS-B", control=list(fnscale=-1))

# Inequality constraints, x1>=0 and x1>=x2
constrOptim(c(3,1), f, g, ui=matrix(c(1,1,0,-1), nrow=2), ci=c(0,0),
            control=list(fnscale=-1))

# Plot the objective function, with inequality constraints
x <- y <- seq(-4, 4, len=50)
z <- outer(x, y, FUN=function(x,y) apply(cbind(x,y), 1, f))
filled.contour(x, y, z, xlab=expression(x[1]), ylab=expression(x[2]),
  color.palette=colorRampPalette(c("blue4","blue3","white","green3","green4")),
  plot.axes={axis(1); axis(2)
             abline(a=0, b=1, lwd=2)
             lines(c(0,0), c(-4,0), lwd=2)
           })
```

# Adaptive Rejection Sampling

---

- Suppose we are interested in generating a random sample from a distribution with density  $f$ . Unfortunately, this distribution is very difficult to sample from. However, there is another density  $g$  that is easy to sample from where for some constant  $c$ ,  $f(x) \leq cg(x)$  for all  $x$ .
- Thus we can generate a sample  $X$  from  $f$  by using rejection sampling,
  1. Generate  $Y$  from  $g$
  2. Generate a random uniform(0,1) number  $U$
  3. If  $U \leq f(Y)/(cg(Y))$  set  $X=Y$
- It may be difficult to determine  $g$ , in this case use adaptive rejection sampling.
- Adaptive rejection sampling (ARS) is a method that performs rejection sampling where the envelope function  $g$  is constructed and refined during the sampling process.
- Method is only applicable to univariate probability density functions that are **log-concave**.

# Adaptive Rejection Sampling

---

- The function `ars()` in the `ars` package performs ARS

```
ars(n=1, f, fprima, ...)
```

<code>n</code>	Sample size
<code>f</code>	Function that returns $\log(f(u, \dots))$ , where $f(u)$ is the density we want to sample from
<code>fprima</code>	$d/du \log(f(u, \dots))$
<code>...</code>	Additional arguments to be passed to <code>f</code>

## Example - Adaptive Rejection Sampling

---

```
library(ars)

# Mixture of two normal distributions with equal weight
f <- function(x) {.5*dnorm(x) + .5*dnorm(x, mean=2)}

integrate(f, -Inf, Inf) # Density integrates to 1
curve(f(x), -3, 5, lwd=2) # Density shape

# Log-density
f.log <- function(x) {log(f(x))}
curve(f.log, -3, 5, lwd=2) # Log-concave down

# Derivative of the log-density
f.log.prime <- function(x) {(-x*.5*dnorm(x) + -(x-2)*.5*dnorm(x, mean=2))/f(x)}

# ARS sample
data <- ars(10000, f.log, f.log.prime)

# Verify
hist(data, freq=FALSE)
curve(f(x), lwd=2, col="blue", add=TRUE)
```