

---

# Basic Statistical Methods

## Lecture 9

---

Nicholas Christian  
BIOST 2094 Spring 2011

# Outline

---

1. Summary Statistics
2. Comparing Means
3. Comparing Variances
4. Comparing Proportions
5. Testing Normality
6. Testing Independence

# Summary Statistics - Vectors

---

## ■ Functions for calculating summary statistics of vector elements

---

<code>mean(x)</code>	Mean of x
<code>median(x)</code>	Median of x
<code>var(x)</code>	Variance of x
<code>sd(x)</code>	Standard deviation of x
<code>cov(x,y)</code>	Covariance of x and y
<code>cor(x,y)</code>	Correlation of x and y
<code>min(x)</code>	Minimum of x
<code>max(x)</code>	Maximum of x
<code>range(x)</code>	Range of x
<code>quantile(x)</code>	Quantiles of x for the given probabilities

---

```
> x <- rnorm(10)
```

```
> y <- runif(10)
```

```
> mean(x)
```

```
[1] 0.2210303
```

```
> cor(x,y)
```

```
[1] 0.3335438
```

# Summary Statistics - Dataframes

---

- Functions for calculating summary statistics of the columns of a dataframe

---

<code>summary()</code>	Summary statistics of each column; type of statistics depends on data type
<code>apply()</code>	Apply a function to each column, works best if all columns are the same data type
<code>tapply()</code>	Divide the data into subsets and apply a function to each subset, returns an array
<code>by()</code>	Similar to <code>tapply()</code> , return an object of class <code>by</code>
<code>ave()</code>	Similar to <code>tapply()</code> , returns a vector the same length as the argument vector
<code>aggregate()</code>	Similar to <code>tapply()</code> , returns a dataframe
<code>sweep()</code>	"Sweep out" a summary statistic from a dataframe, matrix or array

---

- The major difference between `tapply()`, `ave()`, `by()`, and `aggregate()` is the format of the returned object

# tapply()

---

- `tapply()` applies a function to each subset of a vector, where the subsets are determined by the unique combination of levels of factors

```
tapply(X, INDEX, FUN, ..., simplify=TRUE)
```

X	Vector applying FUN to
INDEX	List of factors for determining subsets. Each factor is the same length as X
FUN	The function to apply to each subset
...	Additional arguments to FUN
simplify	If FALSE returns a list, otherwise returns an array

## sweep()

---

- The function `sweep()`, "sweeps out" or subtracts a summary statistic from the specified margins of an array or dataframe.

```
sweep(x, MARGIN, STATS, FUN="-", ...)
```

x	Array or dataframe
MARGIN	Margin to carry out the sweep on
STATS	Array of summary statistics to sweep out
FUN	Function used to carry out the sweep (default is subtraction). Should be a function of two arguments: x and an array with the same dimensions as x generated from STATS
...	Additional arguments to FUN

- The dimension of `STATS` depends on the value of `MARGIN`. So if `MARGIN=2` (columns) and `x` has three columns then `STATS` is a vector of length 3
- `sweep()` appropriately expands `STATS` to match the dimension of `x` and then applies `FUN`.

## Example - Dataframe Summary Statistics

---

```
# Create a dataset of state data from the built-in matrix state.x77
# with two continuous variables (population and income) and
# two factor variables (region and area)
area.factor <- cut(state.x77[, "Area"],
                   breaks=quantile(state.x77[, "Area"], c(0, .25, .75, 1)),
                   labels=c("small", "medium", "large"),
                   include.lowest=TRUE)

state <- data.frame(pop=state.x77[, "Population"],
                   inc=state.x77[, "Income"],
                   area=area.factor,
                   region=state.region)

head(state)
dim(state)
```

## Example - Dataframe Summary Statistics

---

```
# Basic summary statistics for each column
summary(state)

# Variance of first two columns
apply(state[,1:2], 2, var)

# Mean income by region
tapply(state$inc, state$region, FUN=mean)
by(state$inc, state$region, FUN=mean)

ave(state$inc, state$region, FUN=mean)

aggregate(state$inc, list(state$region), FUN=mean)

# Mean population by region and area
with(state, tapply(inc, list(region, area), FUN=mean))
with(state, by(inc, list(region, area), FUN=mean))
with(state, ave(inc, list(region, area), FUN=mean))
with(state, aggregate(inc, list(region, area), FUN=mean))
```

# Return a vector  
# Return a by object,  
# extract elements using [ ]  
# Return a vector, same  
# length as first argument  
# Return a dataframe, subset  
# variable needs to be a list



## Example - Dataframe Summary Statistics

```
# Sweep out the columns means for population and income
# i.e., subtract the overall mean population and overall mean income
# from the population and income for each state, respectively.

# First calculate the statistic we want to sweep out,
# columns means of population and income
col.means <- apply(state[1:5,1:2], 2, mean)
sweep(state[1:5,1:2], 2, col.means, FUN="-")

matrix(rep(col.means, each=5), nrow=5) - state[1:5, 1:2]           # Without sweep()

sweep(state[1:5, 1:2], 2, col.means, FUN=function(x, s)(x-s)/s) # Custom function

# Recall problem 1 part (e) from Homework 1, find  $P(X | Y, Z)$ 
distr <- array(rep(c(.04,.01,.24,.06,.12,.03), 2), dim=c(2,3,2),
               dimnames=list(Y=c(1,-1), X=c(10,20,30), Z=c(0,1)))
marg.YZ <- apply(distr, c(1,3), sum)

# Without using sweep()
expand.marg.YZ <- array(rep(c(marg.YZ), 3), dim=c(2,3,2),
                       dimnames=list(Y=c(1,-1), X=c(10,20,30), Z=c(0,1)))
distr/expand.marg.YZ

# sweep() appropriately expands marg.YZ to match the dimension
# of distr then applies the division function
sweep(distr, c(1,3), marg.YZ, FUN="/")
```

# Apply family and related functions

---

<code>apply()</code>	Apply a function to the margins of an array
<code>lapply()</code>	Apply a function to each element of a list or vector, returns a list
<code>sapply()</code>	Same as <code>lapply()</code> , but returns a vector or matrix by default
<code>vapply()</code>	Similar to <code>sapply()</code> , but has a pre-specified type of return value
<code>mapply()</code>	Multivariate version of <code>sapply()</code> , apply a function to the corresponding elements from multiple lists or vectors
<code>tapply()</code>	Apply a function to each subset of a vector, where the subsets are determined by the unique combination of factor levels
<code>outer()</code>	General outer product of two vectors
<code>sweep()</code>	Sweep out a summary statistic from the margins of an array
<code>replicate()</code>	Repeated evaluation of an expression

---

- These functions allow code to be more compact and clearer

# Contingency Tables

---

- The following functions are used for creating data tables

---

<code>table()</code>	Create a contingency table of the counts at each combination of factor levels
<code>ftable()</code>	Similar to <code>table()</code> , useful for multidimensional tables
<code>xtabs()</code>	Create a contingency table using a formula interface, useful if the data have already been tabulated
<code>prop.table()</code>	Table of proportions, relative to the given margin
<code>addmargins()</code>	Add margins to a table (default is to sum over all margins in the table)

---

- Although it appears that these functions return matrices or arrays, these functions actually return objects of class "table" for `table()`, "ftable" for `ftable()`, and "xtabs" for `xtabs()`. Both `prop.table()` and `addmargins()` return an object of class "table".

## Example - Contingency Tables

---

```
# Create a table of population and area
# Note: prop.table() and addmargins() require a table object
# argument and do not accept factor arguments
pop.area <- with(state.tab, table(pop, area))
addmargins(pop.area)
prop.table(pop.area)                # Overall frequency
prop.table(pop.area, margin=2)      # Column frequency

# Create a multidimension table of population, area, and region
with(state.tab, ftable(pop, area, region))

# To include margins/proportions it is best to add the margins/proportions
# to a table object and then pass the new table object to ftable
pop.area.region <- with(state.tab, table(pop, area, region))
ftable(addmargins(pop.area.region))
ftable(prop.table(pop.area.region))
ftable(prop.table(pop.area.region, margin=3))

# Change row and column variables
with(state.tab, ftable(pop, area, region, row.vars="region"))
with(state.tab, ftable(pop, area, region, row.vars=c("region", "pop")))
with(state.tab, ftable(pop, area, region, row.vars="pop", col.vars="area"))
```

## Example - Contingency Tables

---

```
# Create a dataset where state.tab is already tabulated
# The outputted inc column is actually the frequency of each
# combination since the function we used was length
state.tab.freq <- aggregate(inc~pop+area+region,data=state.tab,FUN=length)
colnames(state.tab.freq)[4] = "freq"

# We also could have used tapply(), but in this case aggregate()
# returns the most useful object
with(state.tab, tapply(inc, list(pop, area, region), FUN=length))

# Since the data is already tabulated xtabs() is the most appropriate
# function. xtabs() is the only table function that has a dataframe
# argument
xtabs(freq~pop+area, data=state.tab.freq)

# If the data was not tabulated omit the left hand side
xtabs(~pop+area, data=state.tab)
```

# Expand a Dataframe

---

- In some cases we may want to expand a dataframe of explanatory variables that has a frequency or count variable into a dataframe with as many repeated rows as specified by a frequency variable. That is, create a dataframe with a separate row for each case.

```
# Create a new dataframe with a separate row for each case
# For example, we want 4 copies of the first row (small, small, Northeast)
state.tab.freq

# Remember a dataframe is like a list where each component of the list
# has the same number of elements
temp <- lapply(state.tab.freq, function(x) rep(x, state.tab.freq$freq))

# Convert to a dataframe and remove the fourth column, freq
as.data.frame(temp)[-4]
```

# Collapse a Dataframe

---

- On the other hand we may want to collapse a dataframe with repeated rows and create a new dataframe with one row for each unique combination of the explanatory variables and a frequency variable giving the number of occurrences.
- There are a couple of ways to do this, use `table()` to include all unique combinations of the variables including those with zero counts and use `aggregate()` to only include the unique variable combinations with non-zero counts.

```
# Collapse state.tab (ignoring inc)
head(state.tab)
```

```
# Include all unique combinations, including those with zero counts
with(state.tab, as.data.frame(table(pop, area, region)))
```

```
# Only include unique combinations with non-zero counts
state.tab.freq <- aggregate(inc~pop+area+region,data=state.tab,FUN=length)
colnames(state.tab.freq)[4] = "freq"
```

# Comparing Means

---

- One sample tests, two sample tests and paired tests of the population mean(s) can be performed using either the  $t$ -test or the Wilcoxon rank-sum test, the non-parametric alternative to the  $t$ -test.
  - ☐ For the  $t$ -test, `t.test()`
  - ☐ For the Wilcoxon rank-sum test, `wilcox.test()`,
- The function `power.t.test()` will compute the power for a  $t$ -test or determine the parameters needed to obtain the desired power



## Example - *t*-test

---

```
# The ToothGrowth dataset looks at the effect of vitamin C on
# tooth growth in guinea pigs. Three dose levels of vitamin C
# (0.5, 1, and 2 mg) were given by two different methods
# (orange juice and ascorbic acid) to 10 guinea pigs.
head(ToothGrowth)

# Only focus on the 1mg dose
Tooth.1mg <- subset(ToothGrowth, dose==1)

# From the boxplot orange juice leads to longer teeth
with(Tooth.1mg, boxplot(len~supp, names=c("Orange Juice", "Ascorbic Acid"),
                        xlab="Supplement Method", ylab="Tooth Length"))

# Two-sample t-test
tt <- t.test(len~supp, data=Tooth.1mg, alternative="two.sided",
             var.equal=FALSE, conf.level=.95)

# Extract results
# See the documentation for a description of the values returned
names(tt)
tt$p.value
tt$conf.int
```

## Example - *t*-test

---

```
# Suppose the dataset does not have an indicator column
uns <- unstack(Tooth.1mg, len~supp)
t.test(uns[,1], uns[,2], alternative="greater", var.equal=TRUE,
       conf.level=.90)

# One-sample t-test just for Orange Juice
# Ho: mu=10 vs Ha: mu!=10
t.test(Tooth.1mg[Tooth.1mg$supp=="OJ", "len"], mu=10)

# Paired t-test
# The sleep dataset is a match paired study with 10 subjects. Compared
# the effect of two drugs on increasing hours of sleep.
t.test(extra~group, data=sleep, paired=TRUE)
with(sleep, t.test(extra[group==1]-extra[group==2]))

# Power
# Exactly one of the parameters n, delta, power, sd, and sig.level must
# be passed as NULL, and that parameter is determined from the others
pw <- power.t.test(delta=1, sd=1, sig.level=.05, power=.8,
                   type="two.sample", alternative="two.sided")

names(pw)
pw$n
```

# Comparing Variances

---

- Use `var.test()` to performs an  $F$ -test to compare the variances of two samples from normal populations.

```
# ToothGrowth dataset where dose=1mg
Tooth.1mg <- subset(ToothGrowth, dose==1)

# Sample Variances
with(Tooth.1mg, tapply(len, supp, var))

# Formula interface
var.test(len~supp, data=Tooth.1mg, alternative="two.sided")

# From histograms the normality assumption of the F-test is not reasonable
par(mfrow=c(1,2))
with(Tooth.1mg, hist(len[supp=="OJ"],
  main="Histogram of length for Orange Juice", xlab="Length"))
with(Tooth.1mg, hist(len[supp=="VC"],
  main="Histogram of length for Ascorbic Acid", xlab="Length"))

# Vector interface,  $H_0: \theta_1/\theta_2 = 2$  vs  $H_a: \theta_1/\theta_2 > 2$ 
uns <- unstack(Tooth.1mg, len~supp)
vt <- var.test(uns[,1], uns[,2], alternative="greater", ratio=2)

names(vt)      # Extract results
vt$p.value
```

# Comparing Proportions

---

- To compare proportions using a normal approximation use `prop.test()`
- To perform an exact binomial test use `binom.test()`
- The function `power.prop.test()` will compute the power for a two-sample test for proportions or determine the parameters needed to obtain the desired power

```
# One-sample test, 17 successes out of 30 trials
prop.test(x=17, n=30)           # Ho: p.hat=.5 vs Ha: p.hat!=.5 (default)
prop.test(x=17, n=30, p=.25)   # Ho: p.hat=.25 vs Ha: p.hat!=.25

# Two-sample test, Ho: p1=p2 vs Ha: p1!=p2 (default)
prop.test(x=c(14,19), n=c(20,30))

# Can also use a matrix to specify the successes and failures
obs <- matrix(c(14,6,19,11), ncol=2, byrow=TRUE,
              dimnames=list(c("Type 1", "Type 2"), c("Success", "Failure")))

# Ho: p1=p2 vs Ha: p1>p2
# Do not use the continuity correction, default does use the correction
pt <- prop.test(obs, alternative="greater", correct=FALSE)

names(pt)
pt$conf.int
```

# Normality

---

- To perform the Shapiro-Wilk test of normality use `shapiro.test()`,  
Ho: Data is normally distributed vs Ha: Data is not normally distributed
- To produce normal quantile-quantile plots use `qqnorm()` along with `qqline()`

```
# Normal Data
```

```
x <- rnorm(100)
```

```
shapiro.test(x)
```

```
qqnorm(x) # Draw Q-Q plot
```

```
qqline(x) # Add a line to the plot that passes through the  
           # first and third quantiles
```

```
# Non-Normal Data
```

```
y <- rgamma(100, shape=2, scale=.5)
```

```
shapiro.test(y)
```

```
qqnorm(y)
```

```
qqline(y)
```

# Independence

---

- To test for correlation/association between paired samples using either Pearson's correlation, Kendall's tau or Spearman's rho, use `cor.test()`
- Kendall's tau or Spearman's rho statistic is used to estimate a rank-based measure of association. These tests may be used if the data do not necessarily come from a bivariate normal distribution.
- All three methods test,  
Ho: No Correlation/association vs Ha: Correlation/association
- To perform a chi-square test of independence use `chisq.test()`
- To perform Fisher's exact test use `fisher.test()`

## Example - Independence

---

```
# Pearson Correlation Test
# Ho: rho=0 vs Ha: rho!=0
x <- rnorm(20)
y <- 3*x + rnorm(20)
plot(x,y)
ct <- cor.test(x, y)
names(ct)
ct$p.value

# Hair and eye color for male and female statistics students
HairEyeColor

# Chi-Square Test for males
# Ho: Independent Ha: Not independent
ct <- chisq.test(HairEyeColor[,,"Male"])

# Warning message since expected count is less than 5
# Note this information is not in the documentation, need to look
# at the body of the chisq.test() function
names(ct)
ct$expected
```